



1. Ablation Study on TabNet’s Sparsity Regularizer


ADULT CENSUS INCOME MASK FIGURES




(a) TabNet ($\lambda_{sparse}=0.0001$, Test Accuracy: 85.46%)



(b) TabNet ($\lambda_{sparse}=0.001$, Test Accuracy: 86.07%)



(c) TabNet ($\lambda_{sparse}=0.01$, Test Accuracy: 86.04%)



(d) TabNet ($\lambda_{sparse}=0.1$, Test Accuracy: 83.41%)

Figure 1: As λ_{sparse} increases, TabNet’s improvement in interpretability is negligible, coupled with a decrease in accuracy.

TabNet (Arik and Pfister, 2020) controls the sparsity of the selected features by employing a sparsity regularizer in the form of entropy (Grandvalet and Bengio, 2004), where $M[i]$ is a learnable mask. Quoted from TabNet,

$$L_{sparse} = \sum_{i=1}^{N_{steps}} \sum_{b=1}^B \sum_{j=1}^D \frac{-M_{b,j}[i] \log(M_{b,j}[i] + \epsilon)}{N_{steps} \cdot B}, \text{ where } \epsilon \text{ is a small number for numerical stability.}$$

Per TabNet’s recommendations, we tuned and produced feature masks in the range of $L_{sparse} = \{0.0001, 0.001, 0.01, 0.1\}$. As depicted in Figure 1, even though we increased TabNet’s sparsity regularizer L_{sparse} , the improvement in interpretability is negligible as almost all the features are reasoned from at each decision step. At the highest L_{sparse} level of 0.1, interpretability might have improved slightly as we start to observe some salient features. However, there is a 2.63% decrease in test accuracy when compared to L_{sparse} of 0.01 (83.41% vs. 86.04%). This proves that naively adding regularization does not work in improving interpretability while maintaining a competitive accuracy.

On the other hand, InterpreTabNet achieves a competitive test accuracy (87.42%) while attaining highly interpretable feature masks (refer to Figure 2 and Table 1 of our main text for InterpreTabNet’s feature mask and accuracy).

References

- S. O. Arik and T. Pfister. Tabnet: Attentive interpretable tabular learning, Feb 2020. URL <https://arxiv.org/abs/1908.07442v4>.
- Y. Grandvalet and Y. Bengio. Semi-supervised learning by entropy minimization. In L. Saul, Y. Weiss, and L. Bottou, editors, *Advances in Neural Information Processing Systems*, volume 17. MIT Press, 2004. URL https://proceedings.neurips.cc/paper_files/paper/2004/file/96f2b50b5d3613adf9c27049b2a888c7-Paper.pdf.